

Concrete strenght prediction using Lasso regretion

```
In [50]: #importing libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Lasso
from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error
import warnings
warnings.filterwarnings('ignore')
```

```
In [10]: df = pd.read_csv('data/concrete_data.csv')

df.head()
```

```
Out[10]:
```

| | cement | blast_furnace_slag | fly_ash | water | superplasticizer | coarse_aggregate | fine_aggregate | age | concrete_cc |
|---|--------|--------------------|---------|-------|------------------|------------------|----------------|-----|-------------|
| 0 | 540.0 | 0.0 | 0.0 | 162.0 | 2.5 | 1040.0 | 676.0 | 28 | |
| 1 | 540.0 | 0.0 | 0.0 | 162.0 | 2.5 | 1055.0 | 676.0 | 28 | |
| 2 | 332.5 | 142.5 | 0.0 | 228.0 | 0.0 | 932.0 | 594.0 | 270 | |
| 3 | 332.5 | 142.5 | 0.0 | 228.0 | 0.0 | 932.0 | 594.0 | 365 | |
| 4 | 198.6 | 132.4 | 0.0 | 192.0 | 0.0 | 978.4 | 825.5 | 360 | |

```
In [11]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1030 entries, 0 to 1029
Data columns (total 9 columns):
 #   Column                                  Non-Null Count  Dtype  
---  -
 0   cement                                1030 non-null   float64
 1   blast_furnace_slag                    1030 non-null   float64
 2   fly_ash                               1030 non-null   float64
 3   water                                 1030 non-null   float64
 4   superplasticizer                      1030 non-null   float64
 5   coarse_aggregate                      1030 non-null   float64
 6   fine_aggregate                        1030 non-null   float64
 7   age                                   1030 non-null   int64  
 8   concrete_compressive_strength         1030 non-null   float64
dtypes: float64(8), int64(1)
memory usage: 72.5 KB
```

```
In [12]: df.shape
```

```
Out[12]: (1030, 9)
```

Data Preprocessing

```
In [13]: #checking for null values
df.isnull().sum()
```

```
Out[13]: cement                                0
blast_furnace_slag                            0
fly_ash                                         0
```

```
water 0
superplasticizer 0
coarse_aggregate 0
fine_aggregate 0
age 0
concrete_compressive_strength 0
dtype: int64
```

```
In [14]: #checking for duplicate rows
df.duplicated().sum()
```

```
Out[14]: 25
```

```
In [21]: df.loc[df.duplicated(), :]
```

```
Out[21]:
```

| | cement | blast_furnace_slag | fly_ash | water | superplasticizer | coarse_aggregate | fine_aggregate | age | concrete |
|------------|--------|--------------------|---------|-------|------------------|------------------|----------------|-----|----------|
| 77 | 425.0 | 106.3 | 0.0 | 153.5 | 16.5 | 852.1 | 887.1 | 3 | |
| 80 | 425.0 | 106.3 | 0.0 | 153.5 | 16.5 | 852.1 | 887.1 | 3 | |
| 86 | 362.6 | 189.0 | 0.0 | 164.9 | 11.6 | 944.7 | 755.8 | 3 | |
| 88 | 362.6 | 189.0 | 0.0 | 164.9 | 11.6 | 944.7 | 755.8 | 3 | |
| 91 | 362.6 | 189.0 | 0.0 | 164.9 | 11.6 | 944.7 | 755.8 | 3 | |
| 100 | 425.0 | 106.3 | 0.0 | 153.5 | 16.5 | 852.1 | 887.1 | 7 | |
| 103 | 425.0 | 106.3 | 0.0 | 153.5 | 16.5 | 852.1 | 887.1 | 7 | |
| 109 | 362.6 | 189.0 | 0.0 | 164.9 | 11.6 | 944.7 | 755.8 | 7 | |
| 111 | 362.6 | 189.0 | 0.0 | 164.9 | 11.6 | 944.7 | 755.8 | 7 | |
| 123 | 425.0 | 106.3 | 0.0 | 153.5 | 16.5 | 852.1 | 887.1 | 28 | |
| 126 | 425.0 | 106.3 | 0.0 | 153.5 | 16.5 | 852.1 | 887.1 | 28 | |
| 132 | 362.6 | 189.0 | 0.0 | 164.9 | 11.6 | 944.7 | 755.8 | 28 | |
| 134 | 362.6 | 189.0 | 0.0 | 164.9 | 11.6 | 944.7 | 755.8 | 28 | |
| 137 | 362.6 | 189.0 | 0.0 | 164.9 | 11.6 | 944.7 | 755.8 | 28 | |
| 146 | 425.0 | 106.3 | 0.0 | 153.5 | 16.5 | 852.1 | 887.1 | 56 | |
| 149 | 425.0 | 106.3 | 0.0 | 153.5 | 16.5 | 852.1 | 887.1 | 56 | |
| 155 | 362.6 | 189.0 | 0.0 | 164.9 | 11.6 | 944.7 | 755.8 | 56 | |
| 157 | 362.6 | 189.0 | 0.0 | 164.9 | 11.6 | 944.7 | 755.8 | 56 | |
| 160 | 362.6 | 189.0 | 0.0 | 164.9 | 11.6 | 944.7 | 755.8 | 56 | |
| 169 | 425.0 | 106.3 | 0.0 | 153.5 | 16.5 | 852.1 | 887.1 | 91 | |
| 172 | 425.0 | 106.3 | 0.0 | 153.5 | 16.5 | 852.1 | 887.1 | 91 | |
| 177 | 362.6 | 189.0 | 0.0 | 164.9 | 11.6 | 944.7 | 755.8 | 91 | |
| 179 | 362.6 | 189.0 | 0.0 | 164.9 | 11.6 | 944.7 | 755.8 | 91 | |
| 182 | 362.6 | 189.0 | 0.0 | 164.9 | 11.6 | 944.7 | 755.8 | 91 | |
| 809 | 252.0 | 0.0 | 0.0 | 185.0 | 0.0 | 1111.0 | 784.0 | 28 | |

```
In [26]: #dropping duplicate rows
df.drop_duplicates(keep='first', inplace=True)
df.shape
```

Out[26]: (1005, 9)

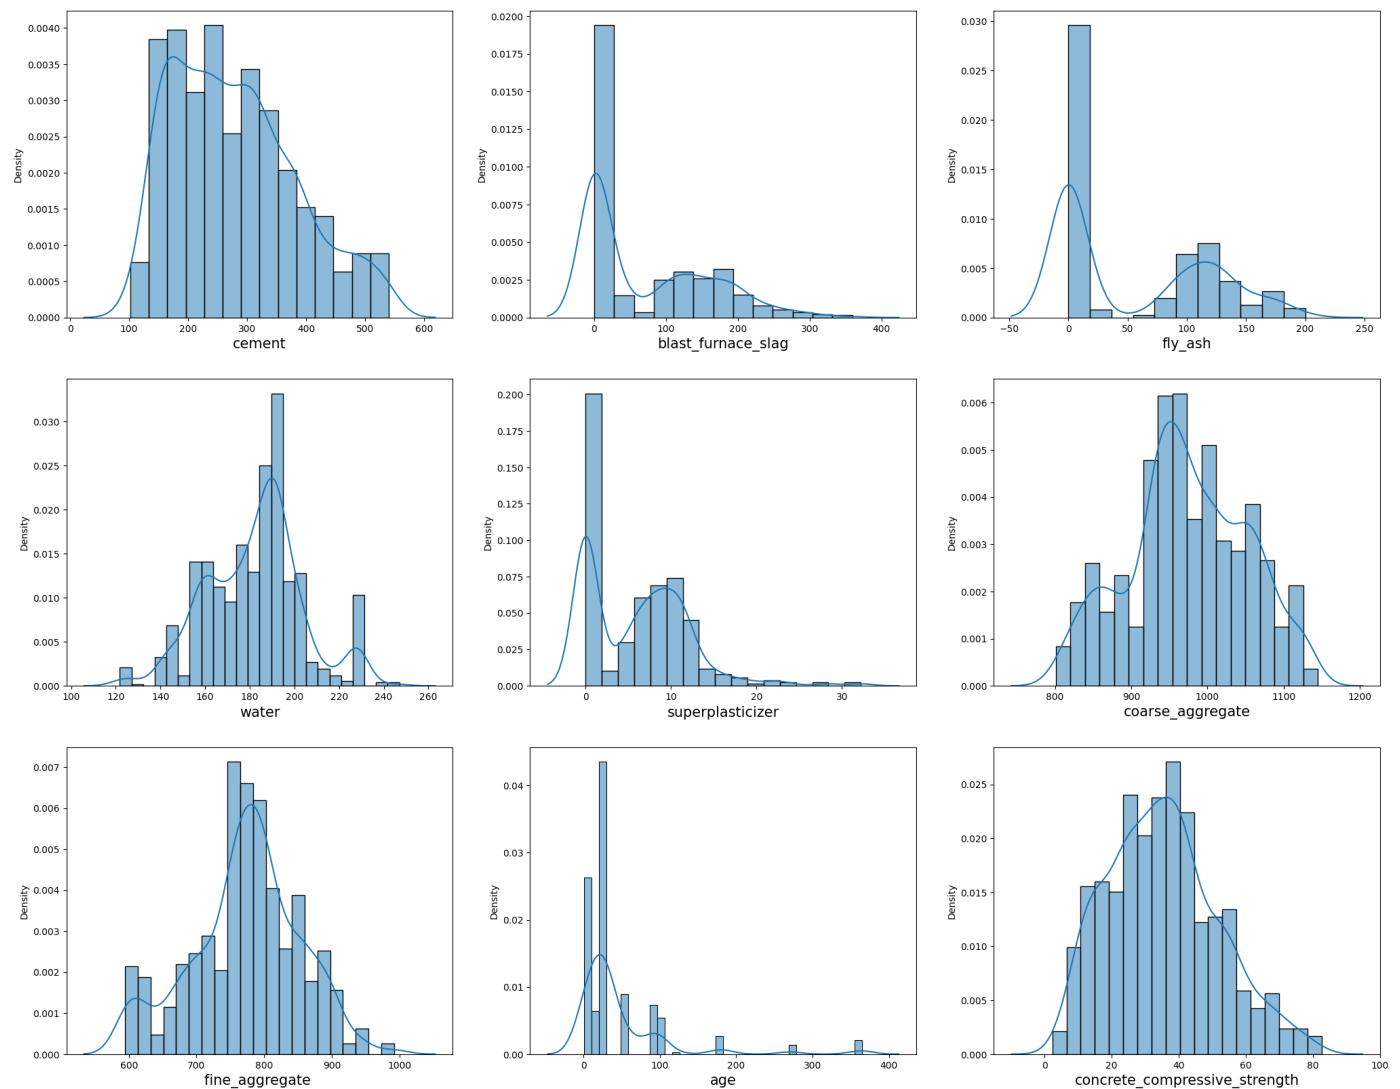
EDA

In [28]: *#plotting distributions of features*

```
plt.figure(figsize = (25, 20))
plotnumber = 1

for col in df.columns:
    if plotnumber <= 9:
        ax = plt.subplot(3, 3, plotnumber)
        sns.histplot(df[col], kde=True, stat='density', kde_kws=dict(cut=3))
        plt.xlabel(col, fontsize = 15)


    plotnumber += 1
```



In [31]: *#correlatin matrix*

```
sns.heatmap(df.corr(), annot=True, vmin=-1, vmax=1)
```

Out[31]: <AxesSubplot: >



- Seems there are no multicollinearity between features

```
In [32]: #pairplot
sns.pairplot(df)
```

```
Out[32]: <seaborn.axisgrid.PairGrid at 0x1e9f8957310>
```



- Cement and strength has linear relationship

Seperating features and target

```
In [33]: X = df.iloc[:, :-1]
y = df.iloc[:, -1]
```

```
In [35]: #splitting dataset into train and test sets
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=1)
```

Feature scaling

```
In [36]: from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
scalerfit = scaler.fit(X_train)
X_train_scl = scalerfit.transform(X_train)
X_test_scl = scalerfit.transform(X_test)
```

Model training

- **### Linear regression**

```
In [56]: lr = LinearRegression()
lr.fit(X_train_scl, y_train)

for i, col in enumerate(X_train.columns):
    print('The coefficient for {} is {}'.format(col, lr.coef_[i]))
```

```
The coefficient for cement is 11.4564804619059
The coefficient for blast_furnace_slag is 7.635182506579105
The coefficient for fly_ash is 5.121438893949539
The coefficient for water is -4.380549966886574
The coefficient for superplasticizer is 1.301366307573757
The coefficient for coarse_aggregate is 0.4455926699073369
The coefficient for fine_aggregate is 0.11771690867432172
The coefficient for age is 7.125608300916846
```

- **### Model evaluation**

```
In [74]: y_pred = lr.predict(X_test_scl)

r2 = r2_score(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)

print('R2 score: ', r2)
print('Mean absolute error: ', mae)
print('Mean squared error: ', mse)
print('Root mean squared error: ', rmse)
```

```
R2 score: 0.5499699178064184
Mean absolute error: 8.917839873008186
Mean squared error: 125.0541420367501
Root mean squared error: 11.182760930859162
```

- **### Lasso regression**

```
In [57]: #finding best value for alpha
from sklearn.model_selection import GridSearchCV

lasso = Lasso()
parameters = {'alpha': [1e-15, 1e-10, 1e-8, 1e-3, 1e-2, 5e-2, 1, 5, 10, 20, 30, 35, 40, 45, 50, 55, 100]}
lasso_regressor = GridSearchCV(lasso, parameters, scoring='neg_mean_squared_error', cv=5)
lasso_regressor.fit(X_train_scl, y_train)

lasso_regressor.best_params_
```

```
Out[57]: {'alpha': 0.05}
```

```
In [58]: lasso_regressor.best_score_
```

Out[58]: -102.87920928665832

```
In [71]: #training the model
lasso = Lasso(alpha=0.05)
lasso.fit(X_train_scl, y_train)

for i, col in enumerate(X_train.columns):
    print('The coefficient for {} is {}'.format(col, lasso.coef_[i]))
```

```
The coefficient for cement is 10.8096250006607
The coefficient for blast_furnace_slag is 6.982465219334084
The coefficient for fly_ash is 4.546702730768315
The coefficient for water is -4.747505308943211
The coefficient for superplasticizer is 1.2699794919238228
The coefficient for coarse_aggregate is 0.0
The coefficient for fine_aggregate is -0.372093231809397
The coefficient for age is 7.034176988616852
```

- **### Model evaluation**

```
In [73]: y_pred = lasso.predict(X_test_scl)

r2 = r2_score(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)

print('R2 score: ', r2)
print('Mean absolute error: ', mae)
print('Mean squared error: ', mse)
print('Root mean squared error: ', rmse)
```

```
R2 score: 0.5492034596064339
Mean absolute error: 8.944123667075363
Mean squared error: 125.26712507143723
Root mean squared error: 11.192279708416745
```

```
In [69]: #coefficients shrinking to 0 when alpha value is high
lasso = Lasso(alpha=3)
lasso.fit(X_train_scl, y_train)

for i, col in enumerate(X_train.columns):
    print('The coefficient for {} is {}'.format(col, lasso.coef_[i]))
```

```
The coefficient for cement is 4.423900784941909
The coefficient for blast_furnace_slag is 0.0
The coefficient for fly_ash is 0.0
The coefficient for water is -0.490293514540079
The coefficient for superplasticizer is 2.6714813724364235
The coefficient for coarse_aggregate is -0.0
The coefficient for fine_aggregate is -0.0
The coefficient for age is 2.78021828050698
```

```
In [ ]: !jupyter nbconvert --to webpdf --allow-chromium-download appliance_energy_prediction.ipynb
```