

# Diamond price prediction with Polynomial regression

```
In [1]: #importing libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import PolynomialFeatures
from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: #reading dataset
df = pd.read_csv('data/diamonds.csv')
df.head()
```

```
Out[2]:
```

	Unnamed: 0	carat	cut	color	clarity	depth	table	price	x	y	z
0	1	0.23	Ideal	E	SI2	61.5	55.0	326	3.95	3.98	2.43
1	2	0.21	Premium	E	SI1	59.8	61.0	326	3.89	3.84	2.31
2	3	0.23	Good	E	VS1	56.9	65.0	327	4.05	4.07	2.31
3	4	0.29	Premium	I	VS2	62.4	58.0	334	4.20	4.23	2.63
4	5	0.31	Good	J	SI2	63.3	58.0	335	4.34	4.35	2.75

```
In [3]: df.shape
```

```
Out[3]: (53940, 11)
```

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 53940 entries, 0 to 53939
Data columns (total 11 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   Unnamed: 0  53940 non-null  int64  
 1   carat       53940 non-null  float64
 2   cut         53940 non-null  object  
 3   color       53940 non-null  object  
 4   clarity     53940 non-null  object  
 5   depth       53940 non-null  float64
 6   table       53940 non-null  float64
 7   price       53940 non-null  int64  
 8   x           53940 non-null  float64
 9   y           53940 non-null  float64
10   z           53940 non-null  float64
dtypes: float64(6), int64(2), object(3)
memory usage: 4.5+ MB
```

## Exploratory data analysis

```
In [5]: #plotting bar charts for categorical variables
```

```
categorical_cols = ['cut', 'color', 'clarity']
```

```
plt.figure(figsize=(15,12))
```



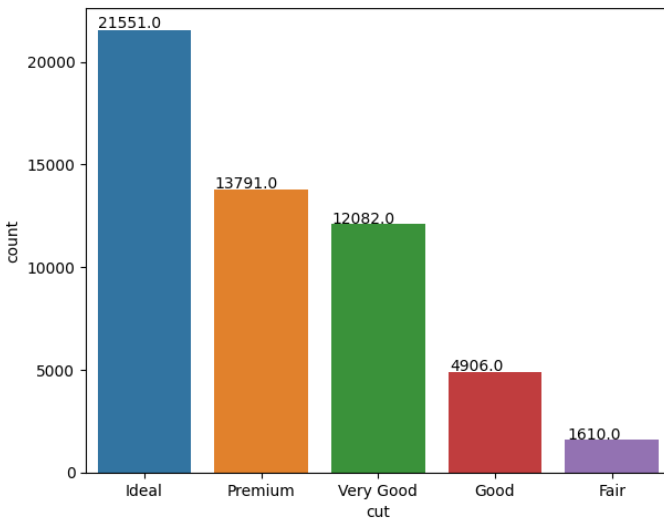
```
for i,col in enumerate(categorical_cols):
```

```
    plt.subplot(2, 2, i+1)
```

```
    ax = sns.countplot(data=df, x=categorical_cols[i], order= df[categorical_cols[i]].va
```

```
    for p in ax.patches:
```

```
        ax.annotate(p.get_height(), (p.get_x() * 1.005, p.get_height() * 1.005))
```



## Inferences:

- Ideal cut diamonds more in numbers while fair cut is the least one
- The worst color 'J' is the rare one, however the bad colors 'G' and 'E' are more in numbers.
- Dimaonds with the best clairty 'IF' and worst clarity 'I1' are less in numbers while 'SI1' and 'VS2' have more number of diamonds.

```
In [6]: #plotting distirbution of numerical features
num_cols = ['carat', 'depth', 'table', 'price', 'x', 'y', 'z']
```

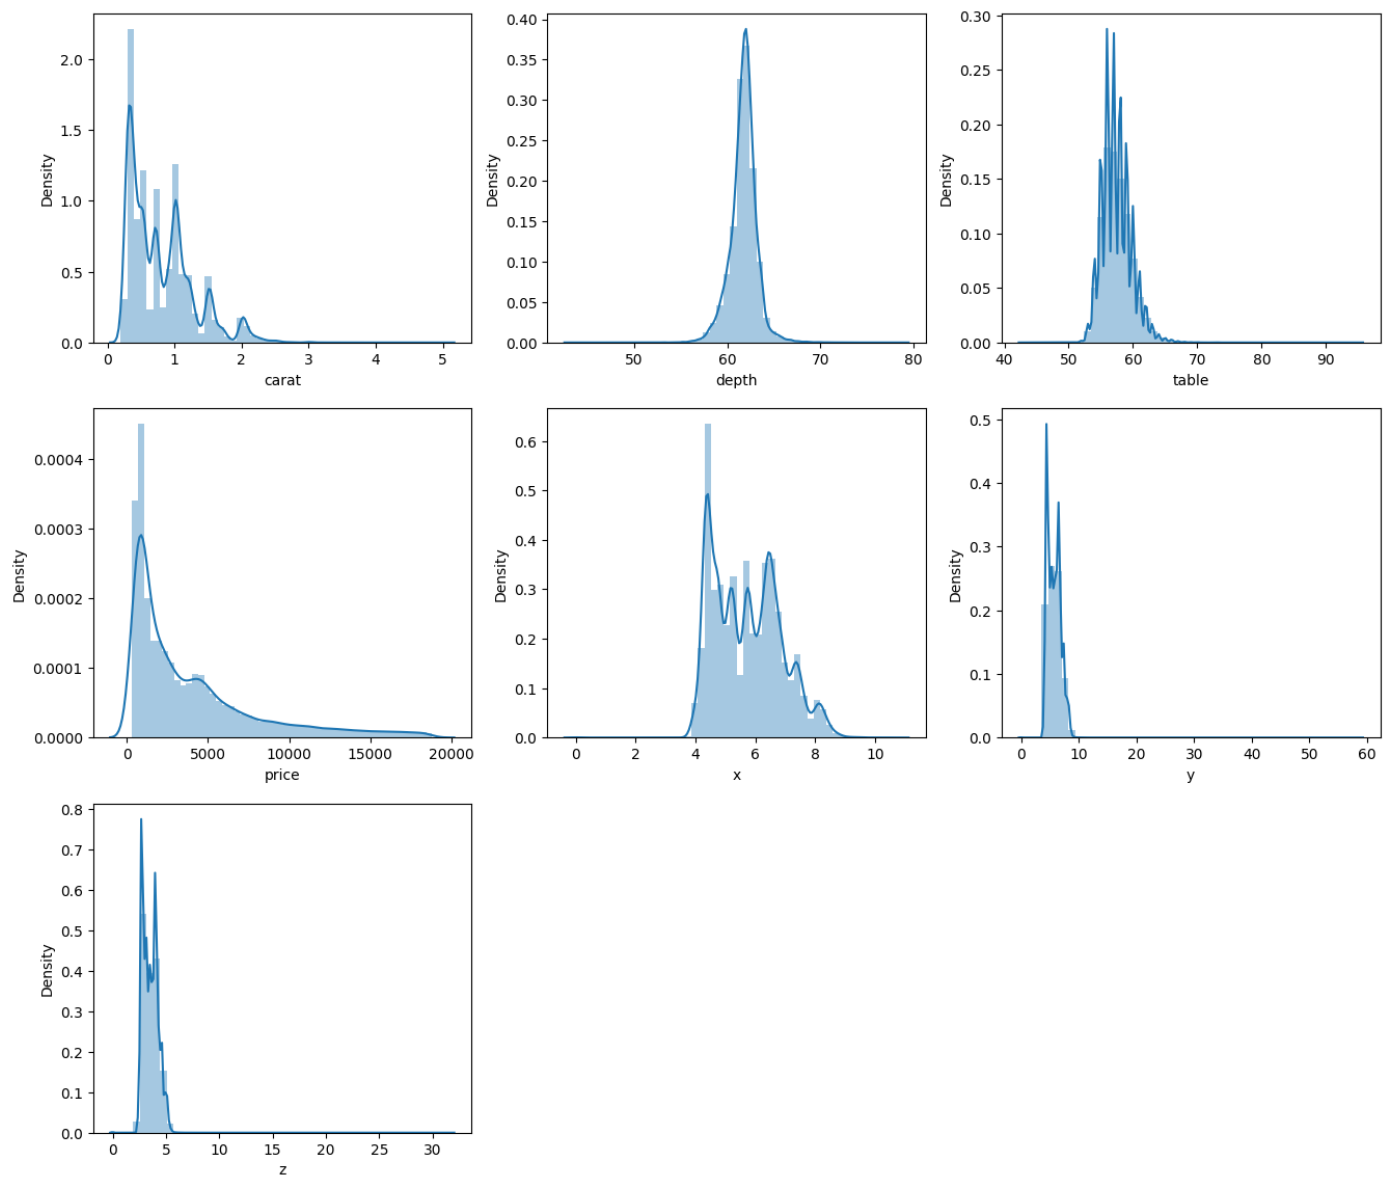
```
plt.figure(figsize=(16,14))
```

```
for i, cols in enumerate(num_cols):
```

```
    plt.subplot(3,3,i+1)
```

```
    #sns.histplot(df[num_cols[i]], kde=True, stat='density', kde_kws=dict(cut=3))
```

```
    sns.distplot(df[num_cols[i]])
```



```
In [7]: df.describe().T
```

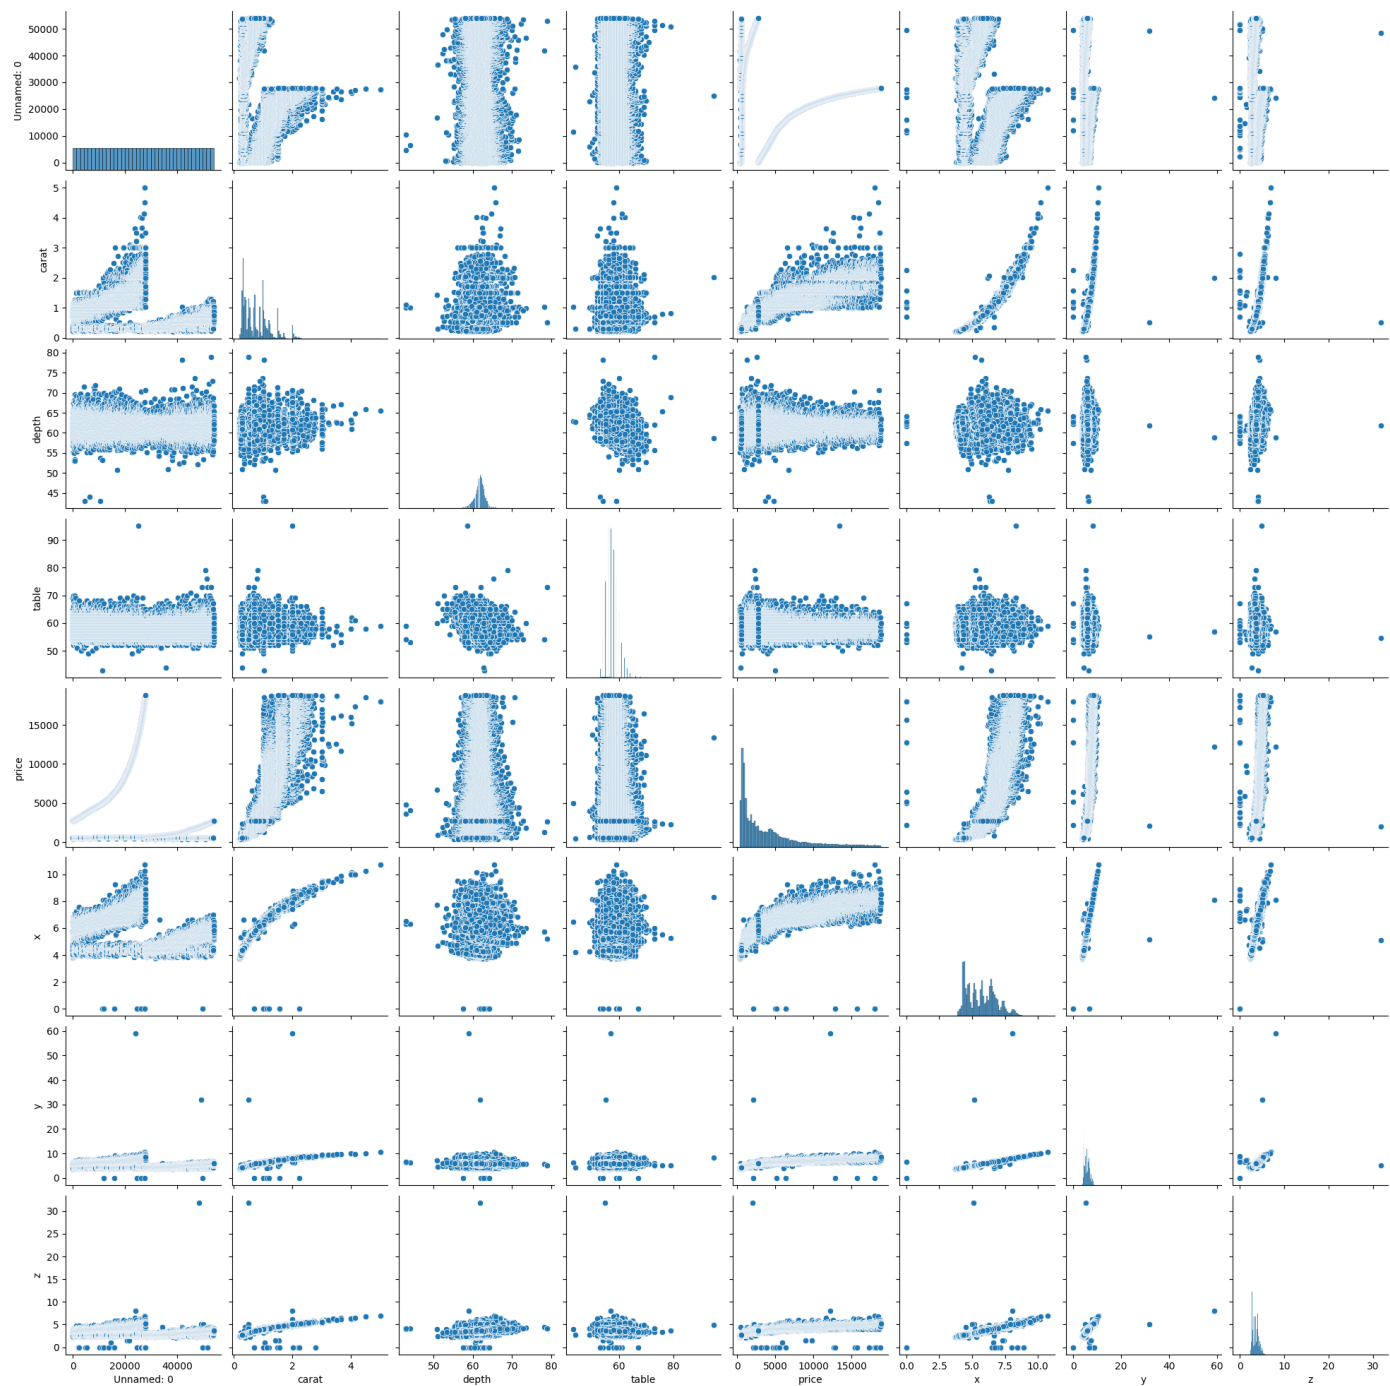
```
Out[7]:
```

	count	mean	std	min	25%	50%	75%	max
<b>Unnamed: 0</b>	53940.0	26970.500000	15571.281097	1.0	13485.75	26970.50	40455.25	53940.00
<b>carat</b>	53940.0	0.797940	0.474011	0.2	0.40	0.70	1.04	5.01
<b>depth</b>	53940.0	61.749405	1.432621	43.0	61.00	61.80	62.50	79.00
<b>table</b>	53940.0	57.457184	2.234491	43.0	56.00	57.00	59.00	95.00
<b>price</b>	53940.0	3932.799722	3989.439738	326.0	950.00	2401.00	5324.25	18823.00
<b>x</b>	53940.0	5.731157	1.121761	0.0	4.71	5.70	6.54	10.74
<b>y</b>	53940.0	5.734526	1.142135	0.0	4.72	5.71	6.54	58.90
<b>z</b>	53940.0	3.538734	0.705699	0.0	2.91	3.53	4.04	31.80

- Min values of x, y and z = 0 implies that there are errors in data (2D diamonds are meaningless)
- Price follows a right skewed distribution

```
In [8]: #pair plot
sns.pairplot(df)
```

```
Out[8]: <seaborn.axisgrid.PairGrid at 0x27d3204bd90>
```



- There is an irrelevant column 'Unnamed: 0'
- columns x,y and z seems to have outliers

## Data preprocessing

```
In [9]: #dropping unwanted columns  
df = df.drop('Unnamed: 0', axis=1)
```

```
In [10]: #checking for null values  
df.isnull().sum()
```

```
Out[10]: carat      0  
         cut        0  
         color     0  
         clarity   0  
         depth     0  
         table     0
```

```
price      0
x          0
y          0
z          0
dtype: int64
```

```
In [11]: #dropping datapoints with x, y and z have min value 0
df = df.drop(df[df['x']==0].index)
df = df.drop(df[df['y']==0].index)
df = df.drop(df[df['z']==0].index)
```

```
In [12]: df.shape
```

```
Out[12]: (53920, 10)
```

```
In [13]: #removing outliers
df = df[(df["depth"]<75)&(df["depth"]>45)]
df = df[(df["table"]<80)&(df["table"]>40)]
df = df[(df["x"]<40)]
df = df[(df["y"]<40)]
df = df[(df["z"]<40)&(df["z"]>2)]
df.shape
```

```
Out[13]: (53909, 10)
```

## encoding categorical variables

```
In [14]: df.head()
```

```
Out[14]:
```

	carat	cut	color	clarity	depth	table	price	x	y	z
0	0.23	Ideal	E	SI2	61.5	55.0	326	3.95	3.98	2.43
1	0.21	Premium	E	SI1	59.8	61.0	326	3.89	3.84	2.31
2	0.23	Good	E	VS1	56.9	65.0	327	4.05	4.07	2.31
3	0.29	Premium	I	VS2	62.4	58.0	334	4.20	4.23	2.63
4	0.31	Good	J	SI2	63.3	58.0	335	4.34	4.35	2.75

```
In [15]: df.cut.unique()
```

```
Out[15]: array(['Ideal', 'Premium', 'Good', 'Very Good', 'Fair'], dtype=object)
```

```
In [16]: df.color.unique()
```

```
Out[16]: array(['E', 'I', 'J', 'H', 'F', 'G', 'D'], dtype=object)
```

```
In [17]: df.clarity.unique()
```

```
Out[17]: array(['SI2', 'SI1', 'VS1', 'VS2', 'VVS2', 'VVS1', 'I1', 'IF'],
              dtype=object)
```

```
In [18]: cut_map = {'Fair':0, 'Good':1, 'Very Good':2, 'Premium':3, 'Ideal':4}
color_map = {'J':0, 'I':1, 'H':2, 'G':3, 'F':4, 'E':5, 'D':6}
clarity_map = {'SI2':0, 'SI1':1, 'VS1':2, 'VS2':3, 'VVS2':4, 'VVS1':5, 'I1':6, 'IF':7}

df['cut_enc'] = df.cut.map(cut_map)
df['color_enc'] = df.color.map(color_map)
df['clarity_enc'] = df.clarity.map(clarity_map)
```


```
In [19]: df.head()
```

```
Out[19]:
```

	carat	cut	color	clarity	depth	table	price	x	y	z	cut_enc	color_enc	clarity_enc
0	0.23	Ideal	E	SI2	61.5	55.0	326	3.95	3.98	2.43	4	5	0
1	0.21	Premium	E	SI1	59.8	61.0	326	3.89	3.84	2.31	3	5	1
2	0.23	Good	E	VS1	56.9	65.0	327	4.05	4.07	2.31	1	5	2
3	0.29	Premium	I	VS2	62.4	58.0	334	4.20	4.23	2.63	3	1	3
4	0.31	Good	J	SI2	63.3	58.0	335	4.34	4.35	2.75	1	0	0

```
In [20]: #correlation matrix
sns.heatmap(df.corr(), annot=True)
```

```
Out[20]: <AxesSubplot: >
```



- price has strong positive correlations with carat, x, y and z
- carat, x, y and z columns have multicollinearity

```
In [21]: #Creating feature matrix and target vector
X = df[['carat', 'depth', 'table', 'cut_enc', 'color_enc', 'clarity_enc']]
y = df['price']
```

```
In [22]: X.shape
```

```
Out[22]: (53909, 6)
```

```
In [23]: #splitting dataset into test and train splits
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=1)
```

# Model training

## 1. Polynomial regression with degree = 2

```
In [24]: #poly regression with degree = 2
poly_reg = PolynomialFeatures(degree=2)
poly_reg = poly_reg.fit(X_train)
X_train_poly = poly_reg.transform(X_train)
X_test_poly = poly_reg.transform(X_test)
```

```
In [25]: X_train_poly.shape
```

```
Out[25]: (37736, 28)
```

```
In [26]: lr = LinearRegression()
lr.fit(X_train_poly, y_train)
```

```
Out[26]: ▼ LinearRegression
LinearRegression()
```

```
In [27]: #model evaluation
y_pred = lr.predict(X_test_poly)

r2 = r2_score(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)

print('R2 score: ', r2)
print('Mean absolute: ', mae)
print('Mean squared error: ', mse)
print('Root mean squared error: ', rmse)
```

```
R2 score:  0.9055883760351068
Mean absolute:  739.7424590347676
Mean squared error:  1456273.0988075286
Root mean squared error:  1206.761409230312
```

## 2. Multiple Linear regression

```
In [30]: mlr = LinearRegression()
mlr.fit(X_train, y_train)
```

```
Out[30]: ▼ LinearRegression
LinearRegression()
```

```
In [31]: #model evaluation
y_pred = mlr.predict(X_test)

r2 = r2_score(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)

print('R2 score: ', r2)
print('Mean absolute error: ', mae)
```

```
print('Mean squared error: ', mse)
print('Root mean squared error: ', rmse)
```

```
R2 score: 0.8783722897802407
Mean absolute error: 884.8190974160511
Mean squared error: 1876073.6763563755
Root mean squared error: 1369.6983888274
```

## Finding best value for degree

```
In [32]: number_degrees = [1,2,3,4,5,6,7]
plt_mean_squared_error = []
for degree in number_degrees:

    poly_model = PolynomialFeatures(degree=degree)


    poly_model = poly_model.fit(X_train)
    poly_X_train = poly_model.transform(X_train)
    poly_X_test = poly_model.transform(X_test)

    regression_model = LinearRegression()
    regression_model.fit(poly_X_train, y_train)
    y_pred = regression_model.predict(poly_X_test)

    plt_mean_squared_error.append(mean_squared_error(y_test, y_pred, squared=False))

plt.scatter(number_degrees, plt_mean_squared_error, color="green")
plt.plot(number_degrees, plt_mean_squared_error, color="red")
```

```
Out[32]: [<matplotlib.lines.Line2D at 0x27d42307cd0>]
```



## 3. Polynomial regression with degree = 5

```
In [ ]:
```

```
In [33]: #poly regression with degree = 5
```



```
poly_reg = PolynomialFeatures(degree=5)
poly_reg = poly_reg.fit(X_train)
X_train_poly = poly_reg.transform(X_train)
X_test_poly = poly_reg.transform(X_test)

plr = LinearRegression()
plr.fit(X_train_poly, y_train)

#model evaluation
y_pred = plr.predict(X_test_poly)

r2 = r2_score(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)

print('R2 score: ', r2)
print('Mean absolute error: ', mae)
print('Mean squared error: ', mse)
print('Root mean squared error: ', rmse)
```

```
R2 score:  0.9561292323588778
Mean absolute error:  468.2426571275253
Mean squared error:  676694.415971052
Root mean squared error:  822.6143786556688
```

## Conclusion

Polynomial regression with degree = 5 gives the best results

```
In [ ]: !jupyter nbconvert --to webpdf --allow-chromium-download diamond_price_prediction.ipynb
```

```
In [ ]:
```