

House price prediction using SVR

```
In [1]: #importing libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.svm import SVR
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
import warnings
warnings.filterwarnings('ignore')
```

Reading Dataset

```
In [2]: df = pd.read_csv('data/USA_Housing.csv')
df.head()
```

```
Out[2]:
```

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price	Address
0	79545.458574	5.682861	7.009188	4.09	23086.800503	1.059034e+06	208 Michael Ferry Apt. 674\nLaurabury, NE 3701...
1	79248.642455	6.002900	6.730821	3.09	40173.072174	1.505891e+06	188 Johnson Views Suite 079\nLake Kathleen, CA...
2	61287.067179	5.865890	8.512727	5.13	36882.159400	1.058988e+06	9127 Elizabeth Stravenue\nDanieltown, WI 06482...
3	63345.240046	7.188236	5.586729	3.26	34310.242831	1.260617e+06	USS Barnett\nFPO AP 44820
4	59982.197226	5.040555	7.839388	4.23	26354.109472	6.309435e+05	USNS Raymond\nFPO AE 09386

```
In [3]: df.shape
```

```
Out[3]: (5000, 7)
```

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 7 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Avg. Area Income                     5000 non-null   float64
1   Avg. Area House Age                  5000 non-null   float64
2   Avg. Area Number of Rooms            5000 non-null   float64
3   Avg. Area Number of Bedrooms         5000 non-null   float64
4   Area Population                      5000 non-null   float64
5   Price                               5000 non-null   float64
```

6 Address 5000 non-null object

dtypes: float64(6), object(1)

memory usage: 273.6+ KB

Exploratory Data Analysis

In [5]:

df.describe().T

Out[5]:







	count	mean	std	min	25%	50%	75%	
Avg. Area Income	5000.0	6.858311e+04	10657.991214	17796.631190	61480.562388	6.880429e+04	7.578334e+04	1.0770
Avg. Area House Age	5000.0	5.977222e+00	0.991456	2.644304	5.322283	5.970429e+00	6.650808e+00	9.5190
Avg. Area Number of Rooms	5000.0	6.987792e+00	1.005833	3.236194	6.299250	7.002902e+00	7.665871e+00	1.0750
Avg. Area Number of Bedrooms	5000.0	3.981330e+00	1.234137	2.000000	3.140000	4.050000e+00	4.490000e+00	6.5000
Area Population	5000.0	3.616352e+04	9925.650114	172.610686	29403.928702	3.619941e+04	4.286129e+04	6.9620
Price	5000.0	1.232073e+06	353117.626581	15938.657923	997577.135049	1.232669e+06	1.471210e+06	2.4690

In [6]:

```
#plotting distributions of features

num_cols = df.columns[:-1]

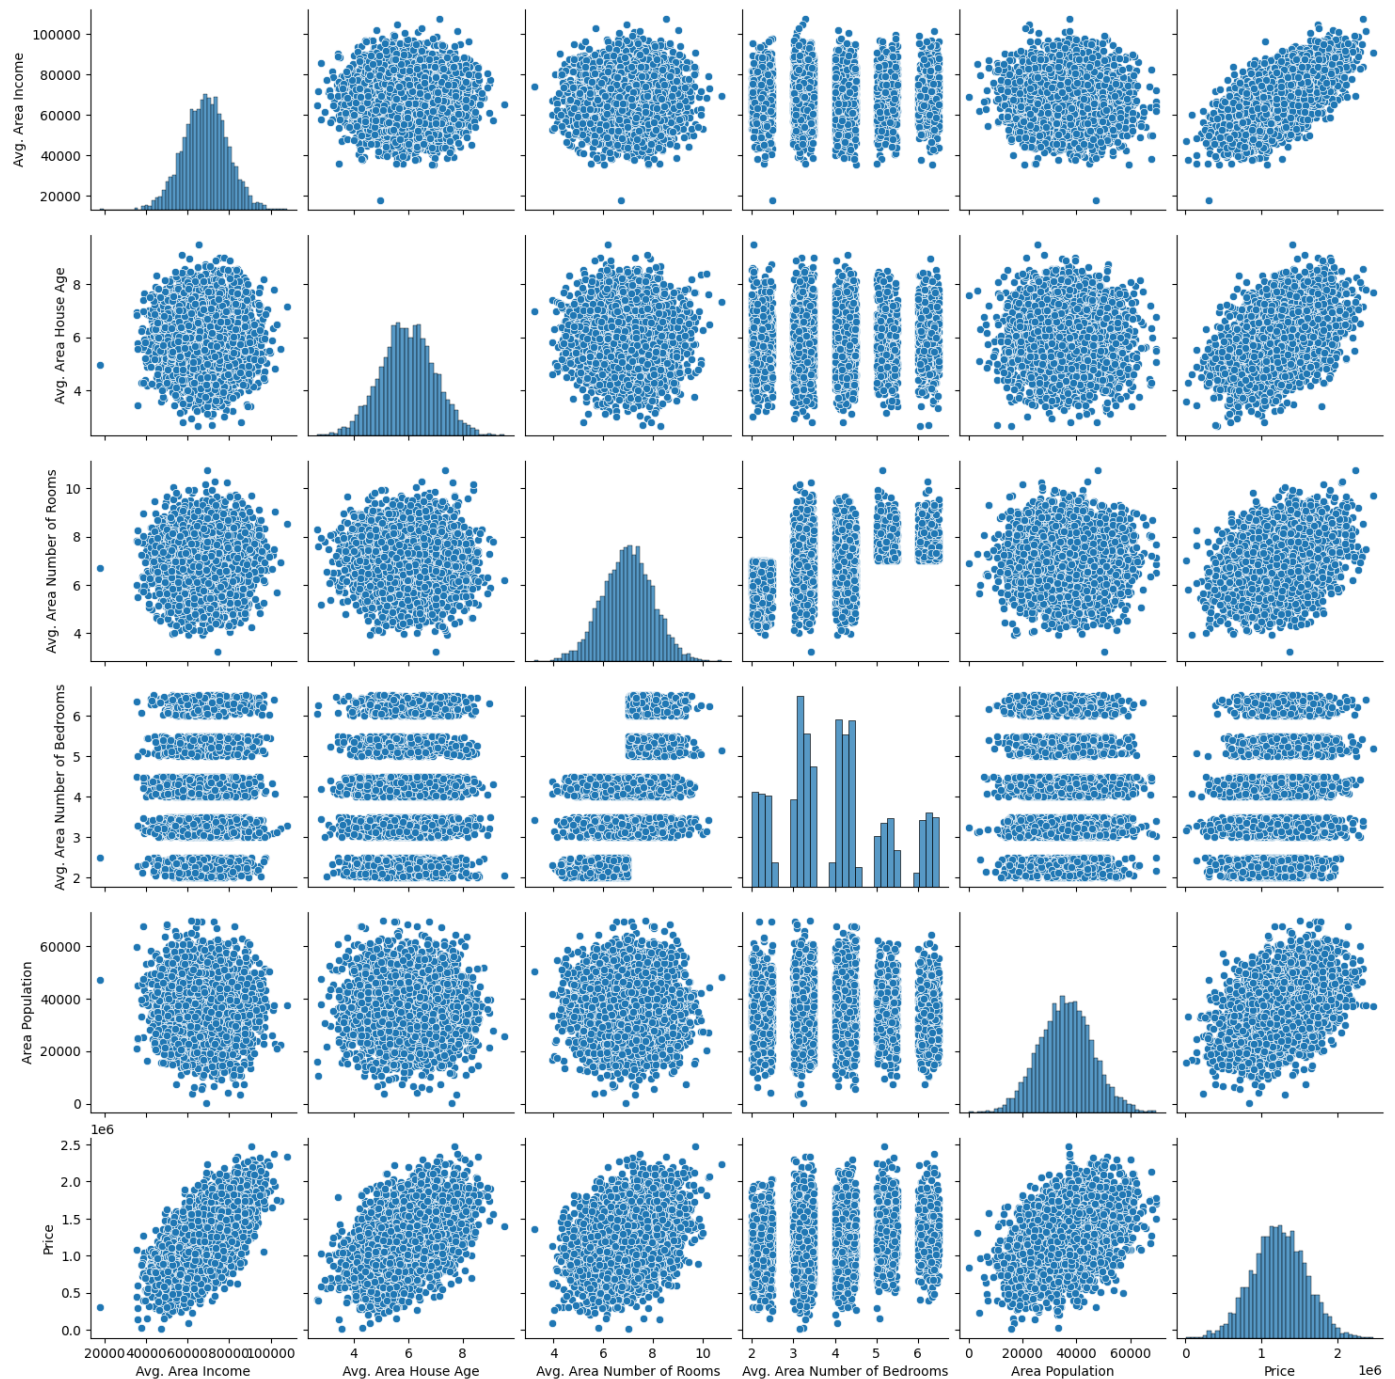
plt.figure(figsize=(20,20))
for i, cols in enumerate(num_cols):
    plt.subplot(3, 2, i+1)
    sns.histplot(df[cols], kde=True, stat='density', kde_kws=dict(cut=3))
    plt.title = cols + 'Distribution'
```



- Features other than Avg. Area Number of Bedrooms follow normal distribution
- Target price also follows normal distribution

```
In [7]: #pairplot
sns.pairplot(df)
```


```
Out[7]: <seaborn.axisgrid.PairGrid at 0x22e2045f430>
```



- The target price has a linear relationship with Avg. Area income
- Avg age, number of rooms and area populations show weak relationship with price

```
In [8]: #correlation heatmap
sns.heatmap(df.corr(), annot=True)
```

```
Out[8]: <AxesSubplot: >
```



- Correlation matrix confirms that Avg. Area income is the one feature that has higher correlation with price
- Avg. age, number of rooms and area populations show weaker correlation
- Avg. Area number of Bedrooms show no to little correlation with price
- There are no multicollinearity between the features
- Avg number of rooms and Avg number of bedrooms show a positive correlation and it is obvious

Data preprocessing

```
In [9]: #checking for null values
df.isnull().sum()
```

```
Out[9]: Avg. Area Income      0
Avg. Area House Age     0
Avg. Area Number of Rooms 0
Avg. Area Number of Bedrooms 0
Area Population          0
Price                   0
Address                 0
dtype: int64
```

```
In [10]: #checking for duplicate values
```

```
df.duplicated().sum()
```

Out[10]: 0

```
In [11]: #dropping Address columns because it seems irrelevent  
df = df.drop('Address', axis=1)
```

```
In [12]: #creating feature matrix and target vector  
X = df.drop('Price', axis=1)  
y = df['Price']
```

```
In [13]: #splitting dataset into train and test sets  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=1)
```

```
In [14]: #feature scaling  
scaler = StandardScaler()  
scaler = scaler.fit(X_train)  
X_train_scl = scaler.transform(X_train)  
X_test_scl = scaler.transform(X_test)
```

Model building

1. Linear regression

```
In [15]: #training the model  
lr = LinearRegression()  
lr.fit(X_train_scl, y_train)
```

Out[15]: ▾ LinearRegression
LinearRegression()

```
In [16]: #model evaluatoin  
y_pred = lr.predict(X_test_scl)  
  
print('r2 score of Linear regression :', r2_score(y_test, y_pred))  
print('MAE of Linear regression :', mean_absolute_error(y_test, y_pred))  
print('MSE score of Linear regression :', mean_squared_error(y_test, y_pred))  
print('RMSE score of Linear regression :', np.sqrt(mean_squared_error(y_test, y_pred)))  
  
r2 score of Linear regression : 0.9215935236936268  
MAE of Linear regression : 82494.73770125784  
MSE score of Linear regression : 10543597313.625992  
RMSE score of Linear regression : 102682.02040097376
```

2. Support Vector Regression

```
In [17]: # #finding best hyperparametes for svr  
# from sklearn.model_selection import GridSearchCV  
  
# params = {'C': [0.1, 1, 10, 100, 1000, 10000, 100000, 1000000],  
#           'gamma': [1, 0.1, 0.01],  
#           'kernel': ['rbf', 'linear', 'poly'],  
#           'epsilon': [0.01, 0.001, 0.0001]}  
  
# svrCV = GridSearchCV(SVR(), params, cv=5, scoring='neg_mean_squared_error')  
  
# svrCV.fit(X_train_scl, y_train)
```

```
In [32]: svr = SVR(kernel='rbf', C=1000000)
svr.fit(X_train_scl, y_train)

y_pred = svr.predict(X_test_scl)

print('r2 score of Linear regression :', r2_score(y_test, y_pred))
print('MAE of Linear regression :', mean_absolute_error(y_test, y_pred))
print('MSE score of Linear regression :', mean_squared_error(y_test, y_pred))
print('RMSE score of Linear regression :', np.sqrt(mean_squared_error(y_test, y_pred)))

r2 score of Linear regression : 0.9109467608056634
MAE of Linear regression : 87242.84408855016
MSE score of Linear regression : 11975305328.990774
RMSE score of Linear regression : 109431.73821607136
```

Conclusion

Linear regression model gives slightly better scores

```
In [ ]: !jupyter nbconvert --to webpdf --allow-chromium-download profit_estimation_of_companies.
```